

The following document describes the test scenarios about possible failures on dataflow job.

The pipeline constructed has the purpose to test the following scenarios:

- 1) An logical error during the execution of some element on the pcollection: Maybe there is an possibility that the constructed pipeline cant process the element received, maybe due the following reasons:
 - a) some kafka message that can not be processed and raise an error
 - b) Maybe there is an external component who change its behavior and causes an unexpected error when some elements are processed, like a call to some external api who responses something not expected, o takes too much time to give the response
- 2) An infrastructure error related to dataflow service, which may cause unexpected behaviors on the pipeline:
 - a) Issues with GCP services related with dataflow: compute engine, cloud storage
 - b) Accidental cancellation of the dataflow job

According to the previous defined scenarios some tests were made on dataflow using the following pipeline:

Read_from_kafka —> test_transformation —> write_to_kafka

For the first scenario we have this test:

- 1) Run the test pipeline with:

```
gcloud dataflow flex-template run "test-kafka-resume-job" \
  --template-file-gcs-location "$TEMPLATE_PATH" \
  --max-workers 2 \
  --region "$REGION" \
  --parameters input_topic="$TOPIC" \
  --parameters output_topic="test-kafka-output-dataflow" \
  --parameters group_id="test-consumer-group" \
  --parameters bootstrap_servers="35.193.114.205:9092" \
  --parameters commit_offset_in_finalize=1 \
  --parameters with_metadata=0 \
  --parameters allow_fail=1 \
  --parameters start_read_time=0 \
  --parameters delay_time=60 \
  --parameters messages_per_delay=10000000 \
  --parameters messages_per_fail=10
```

We can get the following errors on the pipeline:

JOB LOGS

WORKER LOGS

DIAGNOSTICS

Severity

Warning

Filter

Search all fields and values

?

MAX TIME

?

SEVERITY	TIMESTAMP	SUMMARY
>	2023-07-17 18:31:34.102 EDT	message received: 40
>	2023-07-17 18:31:34.102 EDT	making fail message received: 40
ad log entry	2023-07-17 18:31:34.102 EDT	message received after sleep: 40 -- time slept: 0
<div> </div>	2023-07-17 18:31:34.104 EDT	Error processing instruction process_bundle-14-98. Original traceback is Traceback (most recent call last): File "apache_beam/runners/common.py", line 1417, in apache_beam.runners.common.DoFnRunner.process File "apache_beam/runners/common.py", line 837, in apache_beam.runners.common.PerWindowInvoker.invoke_process File "apache_beam/runners/common.py", line 983, in apache_beam.runners.common.PerWindowInvoker._invoke_process File "/template/streaming_kafka_mod.py", line 102, in parse_json_message ValueError: receive message with pair number: 40 During handling of the above exception, another exception occurred: Traceback (most recent call last): File "/usr/local/lib/python3.7/site-

This error is caused intentionally for this test.

To solve this problem without any loose of data, we can update the dataflow job with `allow_fail=0`:

test-
kafka-
resume-
job

test-
kafka-
resume-
job

Streaming

This disable the above error, and we can see how all elements are processed without any error:

Write to kafka

Running

1 sec

1 stage

Logs

HIDE

1

STEP LOGS

encode utf-8.out0

Elements added

300

(Approximate)

Estimated size

33.11 KB

Output Data Freshness

Data freshness by st...

?

Conclusion: No matter the logic code issues or application issues that we can have on the future, if we are able to update the existing running pipeline, no data will be losed or duplicated

Consuming the output topic we can see that all messages has been sent only one time:

```
287 value: 286 count: 1
288 value: 287 count: 1
289 value: 288 count: 1
290 value: 289 count: 1
291 value: 290 count: 1
292 value: 291 count: 1
293 value: 292 count: 1
294 value: 293 count: 1
295 value: 294 count: 1
296 value: 295 count: 1
297 value: 296 count: 1
298 value: 297 count: 1
299 value: 298 count: 1
300 value: 299 count: 1
301
```

For the second scenario we can do:

2) Run the following command

Here we can see that I delete the cluster and the tmp files when not all data was processed an store on the output topic

```
178 value: 287 count: 1
179 value: 288 count: 1
180 value: 289 count: 1
181 value: 290 count: 1
182 value: 291 count: 1
183 value: 292 count: 1
184 value: 293 count: 1
185 value: 294 count: 1
186 value: 295 count: 1
187 value: 296 count: 1
188 value: 297 count: 1
189 value: 298 count: 1
190 value: 299 count: 1
191
```

I leave the process stuck during half hour

```
> ! 2023-07-17 22:58:43.578 EDT adding sleep message received: 240
> ! 2023-07-17 22:58:43.628 EDT The configuration 'group.id' was supplied but isn't a kn...
> ! 2023-07-17 22:58:43.637 EDT The configuration 'group.id' was supplied but isn't a kn...
```

recommended.
[Learn more](#)

Job region ?	us-east1
Worker location ?	us-east1
Current workers ?	0
Latest worker status	Worker pool



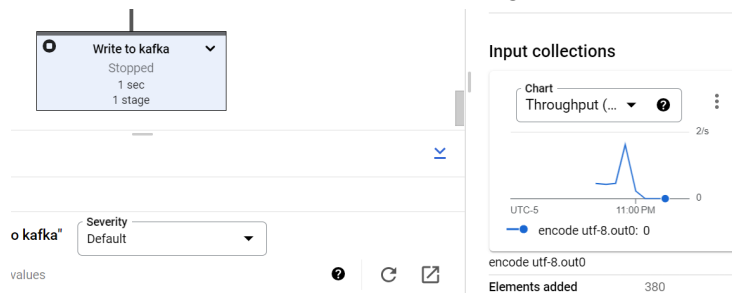
And now I update exactly the same job, no changes

The updated job continue processing the elements that wasn't processed before when I delete the instance group and the gcs staging, It only process the elements that wasn't processed before so no data was lost or duplicated. (300 messages received, counted one time)

```
290 value: 289 count: 1
291 value: 290 count: 1
292 value: 291 count: 1
293 value: 292 count: 1
294 value: 293 count: 1
295 value: 294 count: 1
296 value: 295 count: 1
297 value: 296 count: 1
298 value: 297 count: 1
299 value: 298 count: 1
300 value: 299 count: 1
301
```

Accidental cancelation:

I send another 400 messages and cancel the process just when I receive the first batch of writes on kafka output topic, 80 messages were received:



```
372 value: 691 count: 1
373 value: 692 count: 1
374 value: 693 count: 1
375 value: 694 count: 1
376 value: 695 count: 1
377 value: 696 count: 1
378 value: 697 count: 1
379 value: 698 count: 1
380 value: 699 count: 1
381
```

When I create another job with the same name as the cancelled, it reprocess all the messages, so no checkpoint was shared:

```
373 value: 692 count: 2
374 value: 693 count: 2
375 value: 694 count: 2
376 value: 695 count: 2
377 value: 696 count: 2
378 value: 697 count: 2
379 value: 698 count: 2
380 value: 699 count: 2
```